

Quarterly Business Review

PayFlow Inc. — Q1 2026 | Saravanan Ramesh, Sr. Technical Account Manager

Target SLA	~65ms	\$36/mo	0
99.95% uptime	Avg response time	Monthly infra cost	P1 incidents

Executive Summary

PayFlow Inc. is a Series A FinTech startup (200+ customers, 5,000+ invoices/month, 25% QoQ growth) looking to containerize and deploy on DigitalOcean Kubernetes Service (DOKS). The initial deployment on two s-1vcpu-2gb worker nodes (NYC1) demonstrates production-ready patterns including HPA, cluster autoscaling, pod antiaffinity HA, PostgreSQL persistence, and Kubernetes Secrets management — at \$36/month baseline, 4× cheaper than equivalent AWS EKS architecture.

Current Infrastructure & Performance

Component	Configuration	Cost
DOKS control plane	Managed, NYC1	\$0
Worker nodes (2×)	s-1vcpu-2gb, autoscaler min 2 / max 4	\$24
DO Load Balancer	Standard, Layer 4	\$12
Container Registry	Starter tier, NYC3	\$0
PostgreSQL + PVC	postgres:15-alpine, 1Gi Block Storage	~\$0.10
Total baseline		\$36/mo

Metric	Result	Metric	Result
Avg response time	~65ms	HPA scale-up (2→3 pods)	~60s at 248% CPU
Load test throughput	752 req/sec (50c)	Cluster autoscaler	Node 3 in ~56s
Baseline node CPU	5–6%	Pod recovery time	~22 seconds
vs AWS EKS	4× cheaper (\$36 vs \$159)	P1 incidents	0

Risks & Mitigation

Risk	Severity	Mitigation
No SSL / HTTPS	High	cert-manager + Let's Encrypt — Q2 2026
No user authentication	High	JWT/OAuth2 middleware — Q2 2026
PostgreSQL pod — no HA	High	Upgrade to DO Managed PostgreSQL — Q2 2026
Single-region (NYC1)	High	Active-Passive multi-region NYC1 + SFO3 — Q4 2026
No centralized logging	Medium	Prometheus + Grafana + Loki via Helm — Q3 2026
Manual deployments	Medium	GitHub Actions CI/CD + versioned image tags — Q3 2026
No Layer 7 WAF	Medium	DO LB handles infra DDoS. Add Cloudflare for WAF — Q4 2026

How Containerization Drives Business Outcomes

Performance	Scalability & Reliability	Cost Optimization
• ~65ms avg response time • Load balanced across healthy pods • Readiness probes prevent bad traffic • Rolling deploys with zero downtime • Pod recovery in ~22 seconds	• HPA: 2→3 pods in 60s at 248% CPU • Cluster autoscaler: Node 3 in 56s • Required antiaffinity ensures node HA • Self-healing — failed pods auto-restart • Scales to 1,000+ customers without re-arch	• \$36/mo vs \$159/mo AWS EKS (4× saving) • Free DOKS control plane saves \$72/mo • HPA fills nodes before new Droplets • Single NGINX Ingress LB vs \$12/service • Burst capacity billed hourly, not monthly

Recommendations & Roadmap

Quarter	Priority	Recommendation	Cost
Q2 2026	Critical	Managed PostgreSQL + SSL/TLS + User Auth Replace pod DB with DO Managed PostgreSQL. Add cert-manager + Let's Encrypt. Implement JWT/OAuth2 authentication.	+\$15/mo
Q2 2026	High	Observability stack Prometheus + Grafana + Loki via Helm. App /metrics already Prometheus-ready. DO Alerts + PagerDuty/Slack.	\$0
Q3 2026	Medium	CI/CD + Security + DO Spaces GitHub Actions with versioned image tags. RBAC + Network Policies. DO Spaces for invoice documents. Evaluate node capacity as usage grows.	+\$5/mo
Q4 2026	Low	Multi-region HA + AI expansion + WAF Active-Passive SFO3 for HA. Phase 1 AI invoice chatbot (Claude API via K8s Secret). Cloudflare for Layer 7 WAF.	+\$72-87/mo

High Availability Architecture

	Active-Passive (recommended — Q4 2026)	Active-Active (future — 1,000+ customers)
Traffic	NYC1 serves all; SFO3 on standby	Both regions serve live traffic simultaneously
Failover	Minutes — DNS failover or manual	Instantaneous — automatic LB rerouting
Cost	~\$88/mo (standby runs minimal nodes)	~\$128/mo (both clusters fully provisioned)
DB replication	Asynchronous — simpler to manage	Synchronous — complex, higher latency risk

AI Expansion Roadmap

Phase	Timeline	Description	Cost
Phase 1	Q4 2026	AI invoice chatbot — Claude API via K8s Secret. Stateless pod, scales with HPA. Answers invoice queries, flags overdue accounts, drafts follow-up emails.	API usage only
Phase 2	Q1 2027	Anomaly detection CronJob — scheduled job queries DB, flags unusual payment patterns and churn risk, writes results to alerts table.	~\$0/mo
Phase 3	Future	Self-hosted model on DO GPU Droplet — cost-efficient at high query volume. GPU device plugin integrates with DOKS.	TBD

Account Health	Renewal Risk	Expansion Opportunity
GREEN	Low	High — Managed DB, SSL/Auth, Spaces, Multi-region, AI tooling